


Machine learning II, unsupervised learning and agents: dimensionality reduction



Overview

Motivation

Linear dimensionality reduction (Principal component analysis)

Nonlinear dimensionality reduction (manifold learning)

- Multidimensional scaling (MDS)

- t-SNE

- Others famous methods

- Comments

Motivation

Linear dimensionality reduction (Principal component analysis)

Nonlinear dimensionality reduction (manifold learning)

- Multidimensional scaling (MDS)

- t-SNE

- Others famous methods

- Comments

Dimensionality reduction

We consider the space \mathcal{X} that contains the data, for either supervised or unsupervised learning. In machine learning, we often have $\mathcal{X} \in \mathbb{R}^d$.

- ▶ If d is large (e.g. $\geq 10^4$), the algorithms that run on the data might become too slow to be used, as their algorithmic complexity depends on d (potentially in a quadratic or exponential way, curse of dimensionality)

Dimensionality reduction

We consider the space \mathcal{X} that contains the data, for either supervised or unsupervised learning. In machine learning, we often have $\mathcal{X} \in \mathbb{R}^d$.

- ▶ If d is large (e.g. $\geq 10^4$), the algorithms that run on the data might become too slow to be used, as their algorithmic complexity depends on d (potentially in a quadratic or exponential way, curse of dimensionality)
- ▶ **However**, often the data might actually occupy a **subspace** of lower dimension q , or it may be possible to project the data on such a subspace without losing too much information.
 - ▶ Working in a subspace of lower dimension might speed up the algorithms.
 - ▶ It may also allow visualization of the data.

Main methods of dimensionality reduction

- ▶ **feature selection** : selecting a subset of the original dimensions.
- ▶ **feature extraction** : computing new features from the original features.

Motivation

Linear dimensionality reduction (Principal component analysis)

Nonlinear dimensionality reduction (manifold learning)

- Multidimensional scaling (MDS)

- t-SNE

- Others famous methods

- Comments

`https://scikit-learn.org/stable/modules/unsupervised_reduction.html`

`https://en.wikipedia.org/wiki/Principal_component_analysis`

Introduction

- ▶ Principal Component Analysis is a method used for some unsupervised learning problems.

Introduction

- ▶ Principal Component Analysis is a method used for some unsupervised learning problems.
- ▶ It is based on statistical and geometrical considerations.

Introduction

- ▶ Principal Component Analysis is a method used for some unsupervised learning problems.
- ▶ It is based on statistical and geometrical considerations.
- ▶ It was invented by Pearson at the beginning of XXth century but is still used today.

Introduction

- ▶ Principal Component Analysis is a method used for some unsupervised learning problems.
- ▶ It is based on statistical and geometrical considerations.
- ▶ Applications include :
 - ▶ dimensionality reduction
 - ▶ noise filtering
 - ▶ prediction
 - ▶ general data visualization and analysis

Example

In this paper, astrophysicists use PCA in order to test a new star temperature prediction method [Bermejo et al., 2013]





Figure – PCA used in order to predict temperature.


Problem statement

- We have multidimensional data.



Problem statement


- We have multidimensional data.



We look for the axis that explain or carry the most variations in the data.

Problem statement


- We have multidimensional data.



We look for the axis that explain or carry the most variations in the data.

Thoses axes will be called the **principal components**.

Visual intuition



When looking at this image, what axis would you suggest in order to explain the biggest variation among the data ?

Formalization

We need a mathematical criterion in order to formalize this intuition.

Formalization

We need a **mathematical criterion** in order to formalize this intuition.


In the case of a larger problem with more dimensions, it will not be possible to use visual feedback in order to choose the principal components (the axis).

Furthermore, even in 2D, using only visualization, we can only do a **approximation** of the most relevant axis.


Orthogonal projections

We will restate the problem in a mathematical way, using tools coming from **linear algebra**.
In particular, the concept of **orthogonal projection** is essential.

Orthogonal projection



Orthogonal projection



Exercice 1 : Mathematical problem Can you link orthogonal projections to the problem of finding the most relevant axis ?

Inertia

- ▶ The **inertia** related to an axis will be a measure of the relevance of an axis.
- ▶ If x_i is a sample taken from n samples, and Δ an axis, we call $p_{x_i \rightarrow \Delta}$ the orthogonal projection of x_i on the axis Δ .

Inertia

- ▶ The **inertia** related to an axis will be a measure of the quality of an axis.
- ▶ If x_i is a sample from n samples, and Δ an axis, we call $p_{x_i \rightarrow \Delta}$ the orthogonal projection of x_i on the axis Δ .
- ▶ The inertia related to Δ is :


$$I_{\Delta} = \frac{1}{n} \sum_{i=1}^n d^2(x_i, p_{x_i \rightarrow \Delta}) \quad (1)$$

Inertia


Exercise 2: No inertia

In what situations could we have $I_{\Delta} = 0$?


Inertia



Inertia



Inertia



Inertia

Exercice 3 : Computing an inertia

Please `cd pca/custom data/` and use the file `inertia.py` in order to represent the projections of the data onto a chosen axis and to compute the inertia.

What is your optimal axis ?

Remark

Minimizing I_{Δ} is the same as maximizing I_{Δ^*} where Δ^* is the supplementary orthogonal space.

Several principal components

- ▶ In $2D$ (like in the former example), we can only have 1 or 2 principal component.

Several principal components

- ▶ In $2D$ (like in the former example), we can only have 1 or 2 principal component.
- ▶ However, when the data have more dimensions, it is possible to have **several principal components** (as many as the dimensionality of the data).
- ▶ the data are then projected on these components.

Link with expected values

We will connect PCA and the inertia to **expected values and variances**.

Expected value (espérance)

- ▶ For a discrete random variable X that takes the values x_i with probability p_i :

$$E(X) = \sum_{i=1}^n p_i x_i \quad (2)$$

- ▶ For a continuous random variable X with density $p(x)$:

$$E(X) = \int x p(x) dx \quad (3)$$

Expected value (espérance)

Exercice 4 : Computing an expected value

- ▶ For a discrete random variable X that takes the values x_i with probability p_i :

$$E(X) = \sum_{i=1}^n p_i x_i \quad (4)$$

- ▶ For a continuous random variable X with density $p(x)$:

$$E(X) = \int x p(x) dx \quad (5)$$

Compute the expected value of the dice game.

Variance

$$\text{var}(X) = E\left((X - E(X))^2\right) \quad (6)$$

Variance and Covariance

$$\text{var}(X) = E\left((X - E(X))^2\right) \quad (7)$$

$$\text{cov}(X, Y) = E\left((X - E(X))(Y - E(Y))\right) \quad (8)$$

Probability vs Statistics

(Advanced notions)

- ▶ from a dataset of samples (x_1, \dots, x_n) , we can only compute **statistics**
- ▶ for instance
 - ▶ sample mean :

$$\hat{x} = \frac{1}{n} \sum_{i=1}^n x_i \quad (9)$$

- ▶ unbiased sample variance :

$$\frac{1}{n-1} \sum_{i=1}^n (x_i - \hat{x})^2 \quad (10)$$

The $\frac{1}{n-1}$ factor instead of $\frac{1}{n}$ requires some theory to understand.

Variance and Inertia

Exercise 5 : **Linking variance and inertia**

What is the relationship between variance and inertia ?

Optimization

- ▶ In real applications, algorithms or analytic solutions are used in order to find the optimal axes.

Optimization with analytic solution

- ▶ Methods include the **Lagrange multiplier method** (optimisation method)
- ▶ The method shows that we need to compute the **eigenvectors** of the **covariance matrix**.
- ▶ See also : Gram matrix

Optimization with algorithms

- ▶ One can also use an approximation method to find the first principal component, called the **power iteration** algorithm.
- ▶ It is useful when the dimensionality is high, when the Lagrange multiplier method becomes too slow, since it involves computation of the covariance matrix.

Optimization with algorithms

Exercise 6 : Complexity of computing the covariance matrix. If n is the number of datapoints, and p the number of variables, what is the complexity of this calculation ?

Optimization with algorithms

Exercise 6 : Complexity of computing the covariance matrix.

If n is the number of datapoints, and p the number of variables, what is the complexity of this calculation ?

The power iteration algorithm only involves cnp operations, where c is a constant way smaller than p , so it is way faster.

PCA with scikit-learn

- ▶ We will use scikit-learn in order to perform PCA on our dataset.
- ▶ <https://scikit-learn.org/stable/modules/decomposition.html>
- ▶ <https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html>

PCA with scikit-learn

Exercise 7 : scikit-learn

Please use the file `pca_sklearn.py` in order to find the principal components on the dataset of the previous exercise.

Applications

- ▶ Let us review some applications of the method.

Iris dataset

- ▶ We can perform PCA on the iris dataset as in the file **iris/pca-iris.py**.

Iris dataset





Figure – PCA performed on the iris dataset. We see that the principal components are able to separate the data.

PCA on digits

- ▶ We will also perform the PCA on a dataset consisting in 8×8 pixels images of digits.
- ▶ The idea is to see if the PCA can help us visualize structure in the data.




PCA on digits

Exercise 8: Performing PCA

Please use the file `pca_digits.py` in order to apply PCA to this dataset.

PCA on digits



Number of components


What is a relevant number of components for PCA ?

Number of components

Exercise 9 : Choosing the number of components

Use the file `pca_digits_variance.py` in order to determine how many components are necessary in order to keep 75% of the variance in the digits dataset.

Number of components




Number of components

Exercise 10: Redundancy

What happens with the data contained in **redundancy/redundant_data.npy**? You can analyze them with the file **redundancy/pca_variance.py**.

Number of components



Number of components

Conclusion : PCA can help determine whether some components carry no information in the data.

Shortcomings of PCA

PCA is sensitive to :

- ▶ outliers
- ▶ initial data scaling

Shortcomings of PCA

PCA is sensitive to :

- ▶ outliers
- ▶ initial data scaling
- ▶ Many variants and heuristics exist on this topic.

Asset of PCA

Reducing the dimensionality of the data might be very helpful in a situation where you need to train a classification algorithm on the projected data. The algorithm might be way faster on data that are in a smaller space. However, a sufficient amount of information should be kept during the reduction, hence the necessity of heuristics.

See also

- ▶ kernel-PCA
- ▶ Randomized PCA

Motivation

Linear dimensionality reduction (Principal component analysis)

Nonlinear dimensionality reduction (manifold learning)

- Multidimensional scaling (MDS)

- t-SNE

- Others famous methods

- Comments

<https://scikit-learn.org/stable/modules/manifold.html>
[https://en.wikipedia.org/wiki/Nonlinear_
dimensionality_reduction](https://en.wikipedia.org/wiki/Nonlinear_dimensionality_reduction)

Manifolds

Fundamental hypothesis (as before) :

In some cases, although the dimensionality of the data is high (as for images), the data lie on a specific subpart of the linear space containing them.

This "subpart" is assumed to be a **manifold** (variété).

Manifolds




Figure – 1 dimensional manifold, in \mathbb{R}^2

Linearity

- ▶ In the case of PCA, the manifolds are **linear subspaces** of the ambient space : PCA is a **linear method** and does **linear projections**.
- ▶ A linear mapping in a linear space E is a mapping u such that

$$\forall x, y \in E, u(x + y) = u(x) + u(y) \quad (11)$$

Nonlinearity

- ▶ However, in many situations, the data might lie on **nonlinear** manifolds.
- ▶ **Manifold learning** is the unsupervised meaning of these manifolds in order to study the structure of the data.

- └ Nonlinear dimensionality reduction (manifold learning)
 - └ Multidimensional scaling (MDS)

Multidimensional scaling (MDS) / Positionnement multidimensionnel


- Projects the data in a smaller subspace trying to preserve pairwise similarities between points.

t-SNE

- ▶ t-Stochastic Neighbor Embedding
[Van Der Maaten and Hinton, 2008]
- ▶ <https://lvdmaaten.github.io/tsne/>

- └ Nonlinear dimensionality reduction (manifold learning)
- └ t-SNE

t-SNE on MNIST



- └ Nonlinear dimensionality reduction (manifold learning)
- └ Others famous methods

See also



- ▶ Isomap
- ▶ LLE

Comments

Some disadvantages of non linear manifold methods :

- ▶ Hard to determine a good output dimension (whereas in PCA we can use explained variance) and it is hard to interpret the embedded dimensions (whereas in PCA we know what they mean).
- ▶ Depends on the number of neighbors chosen (if relevant)
- ▶ Often computationally slower.

References

-  Bermejo, J. M., Ramos, A. A., and Prieto, C. A. (2013). Astrophysics A PCA approach to stellar effective temperatures. 95 :1–9.
-  Van Der Maaten, L. J. P. and Hinton, G. E. (2008). Visualizing high-dimensional data using t-sne. *Journal of Machine Learning Research*, 9 :2579–2605.